The Benevolent Bandwidth Foundation
Charter and Operating Principles

The Benevolent Bandwidth foundation (b2) is a nonprofit organization that aims to create a global open-source community for good. Benevolent Bandwidth's mission is to solve problems with transparent and open technologies that deliver public benefit by prioritizing safety, fairness, and trust, without collecting any human data. The foundation envisions building a world where the benefits of AI are shared by everyone, strengthening fairness, opportunity, and human agency.

We have created this non-profit to build a different kind of tech, create a different kind of AI, and a different kind of an organization. Our goal is to reduce barriers to using AI for all humanity (serving the many, not the few), reduce barriers to solving tough human problems, and reduce barriers to open source contribution and giving back. We are hyper-optimized to solve problems but not optimized for user engagement. As such, we will never distract humans while they are using one of our apps with ads, personalization that serves no purpose, gamification, or irrelevant content just to keep users longer in the application.

The main objectives of b2 can be articulated as:

- Identify high leverage problems and ideas for humanity that can be solved with technology
- Identify, motivate and onboard technical advisors and contributors
- Connect projects with contributors
- Execute on projects, deliver them, communicate wins, promote adoption
- Maintain and support live projects

The foundation will live by and adhere to a set of principles. These can be articulated as:

- We start with the problem not the solution or technology and build useful tools that are practical, usable, and easy to adopt
- We use AI for public benefit, prioritizing safety over speed, focus on fairness, and real-world positive impact on humans
- We are open source by default - code, docs, and decisions are transparent
- We never retain human data (user / behavior / click / personal data), no tracking, content is never logged, we never sell or share data, and we never render ads in our tools
- We do not support any political or ideological agendas, political parties or candidates, or partisan positions. We do not serve the mission or act on behalf of specific organizations, governments, countries, or institutions. Funding does not buy influence over b2's decisions or priorities
- We are humble, love learning and are genuinely curious. We invite feedback, own mistakes, and pivot when evidence points to a new direction.
- We are an open and inclusive community. We are open to anyone who wants to contribute. We provide support and stand on each other's shoulders
- We source ideas from anywhere (e.g., foundation members, non-profits, NGOs, etc.) as long as it adheres to the principles above

The b2 foundation does not operate as a “traditional organization” with hierarchies. No one is above anyone else. We are all equals, regardless of how we contribute. We stand on each other’s shoulders, support each other and work towards the singular goal of helping humanity. The foundation will, however, have three groups of contribution opportunities.

Benevolent Board of Directors: These individuals are elected officers of the foundation, legally required by Commonwealth of Massachusetts law (where the foundation is registered). Directors are individuals who are committed to the mission and principles of the Foundation and have fiduciary responsibilities. Each director shall serve a term of two (2) years and may be re-elected for successive terms without limit. They are volunteers, just like everyone else, and are not compensated for their services. Directors have fiduciary and legal authority over the Foundation. Directors are a subset of the Foundation Advisors group. Every Director is, by definition, also a Foundation Advisor.

Executive Committee: These individuals are volunteers who administer the foundation. Specific functions within this committee include: TA recruiting and onboarding, and engagement. Contributor recruiting (university partnerships) and onboarding. TA + Contributor engagement and benefits/incentives. Nonprofits/NGOs partnerships, roundtables, and idea pipeline. Project progress monitoring and product adoption. AI safety, ethics, security, guardrails and governance. Legal matters. B2 tech stack and infrastructure. Product maintenance and continuity plans. Fundraising, accounting, finance. PR, comms, social media, conference relations. Career coaching, resume review, mock interview

Benevolent Foundation Advisors (FAs): These individuals are the stewards of the foundation as a whole, review the project portfolio, identify new ideas, connect the foundation with potential contributors at scale. They are focused on ensuring the organization is living by its mission, vision, and principles. These individuals are globally recognized and accomplished leaders from industry, academia, non-profit, or governmental agencies. Their names will carry weight, open doors and unblock the organization. At scale, we expect Technical Advisors to serve as FAs as they contribute, deliver and gain leadership experience.

Benevolent Technical Advisors (TAs): These individuals own projects, oversee design, approach, test, and release contributions. They identify new ideas, connect potential contributors with the organization, own github repos, and develop contributors. They are highly technical individuals with many years of experience in tech, data, AI, product, design, and leadership. Their names and backgrounds will create excitement among contributors and be one of the main reasons contributors join the foundation. At scale, we expect contributors to serve as TAs as they grow, deliver and gain leadership experience.

Benevolent Contributors (Open Source Fellow): These individuals will code, test, document, and PR/commit high quality new features. They design, build UI/UX flows, develop go to market strategies. They drive the projects to completion. These folks include anyone and everyone. They are, however, generally people who are at the tail end of their studies (undergraduate or graduate) or early in their careers.

We believe it is imperative that everyone who is part of the foundation feel like they belong and their basic human needs such as being recognized, rewarded, validated, bringing their full self, getting their voices heard, etc. are met. We are building a **community** of contributors that are motivated by doing good and contributing, who are mission driven, and are willing to learn. To support this we will establish the following ceremonies, activities and events:

Ceremonies:

- Bi-monthly Board meeting: This meeting is a legally required meeting where certain decisions are made (please see Board of Directors decision boundaries section below).
- Monthly FA roundtable: The goal is to review and discuss backlog, direction, alignment with mission, vision and principles, contributor engagement, budget, etc.
- Monthly TA roundtable: The goal is to review and discuss progress, blockers and needs of our technical advisors.
- Monthly community demos: The goal is to review demos of tools, progress, and blockers. Open to all b2 members.
- Weekly / bi-weekly TAs & collaborators: The goal is to review PRs, approach, design, and joint problem solving.

Community building events:

- Local, in person, social meet ups between TAs and/or collaborators. In collaboration with other open source communities / organizations.
- In person annual planning for FAs.

Recognition and Rewards:

- Contributor of the month
- Most impactful project of the quarter
- Personal milestone celebrations

Our core technical/build principles mirror our foundation principles. Specifically, these include:

- We aren't just building products; we’re building a culture of responsible, high-velocity engineering. Our main differentiators are **quality, safety, privacy, transparency, and trust.** We shall never compromise those in favor of getting something out quickly.
- **Speed still matters:** We do want our projects to see daylight and not get stuck in perpetual build mode. The best way to ensure we are both safe and fast is to set the phases as simple but useful as possible. Thinking through the deliverable and being deliberate before jumping to build is critical.
- **AI Everywhere:** We use AI-assisted tools for coding, design, planning, and documentation. It’s our force multiplier. We double check outputs, test rigorously, and validate before releasing work coming out of AI.
- **No human data.** We will never own, store, or persist human data. We define human data as PII (name, address, phone, email, IP address), behavioral & interaction data (clicks, inputs), or any other data that is generated by a human. We consider model

embeddings of any human data also as human data. We can collect machine data (e.g., server uptimes, number of runs, number of users, etc.) to ensure we are solving the problems we set out to solve. We must document what we are collecting, including clearly identifying how it is not "human data".

- **Safety guardrails**: Our systems must handle failures gracefully. Use of AI must have guardrails by default and, if not, a clear articulation as to why we don't need them. SLAs are well articulated and published when relevant. HITL is embedded when needed.
- **Laws and regulations.** We comply with all laws and regulations without cutting corners.
- We always evaluate **smaller, optimized models/distilled LLMs** first to reduce our environmental impact.
- **Agentic Frameworks:** We use cutting-edge agentic approaches, but do so with a "mechanics-first" mindset. We build with a deep understanding of the underlying logic, not just black-box implementations. In many cases, we opt to build our own agentic orchestration vs. using abstraction tools.
- **Vertical Integration:** We generally avoid building generic or general purpose tools. We don't build a "legal assistant" but a "California Real Estate Compliance Auditor for Multi-Family Units." We solve specific, high-value problems deeply.
- **Open source everything:** Docs, architectural decisions, models, and frameworks, are open-sourced. Transparency is our advantage; we have nothing to hide.
- **Public Roadmaps:** We publish our product roadmaps publicly via issues in Github.
- **Standardized PRDs:** We create Product Requirement Documents (PRD) for all phases of a product and ensure they are open sourced in our repo.
- **Verification & Fact-Checking:** If an AI generates an output in our tools we cite the source (where applicable/possible).
- **System Controls:** We build robust guardrails, continuous monitoring, a "kill switch," and human-in-the-loop processes from phase one.

Frequently Asked Questions (FAQs):

- **Is this organization registered as a nonprofit?**
  - Yes. As of 3/9/2026 b2 is registered in the Commonwealth of MA. The Certificate of Organization, Board meeting minutes and By-laws can be found here.
- **How do we source projects / new ideas?**
  - Sources of projects or ideas can be foundation members, nonprofit organizations, NGOs, local governments, etc. as long as the ideas adhere to b2's principles.
- **Will we raise funds?**
  - Currently this initiative is self funded. Once we complete the next phase (4-5 projects) we will revisit this need and may opt to collect donations, grants, etc. In order to do this we will need to establish as a formal nonprofit.
- **What is the path for growth?**
  - MVP: Research agent - phase completed.
  - Phase-1: 4-5 projects. 4-5 TAs. 15-20 contributors - phase in progress.
  - Phase-2: 15-20 projects. 10-15 TAs. 80-100 contributors.
- **How do we prioritize / decide on what to work on?**
  - We will use the following rubric as a starting point:

        - Can the project be open sourced - hard line?
        - Does the tool require user data collection and storage - hard line?
        - What is the time to deliver and impact?
        - What is the size of impact: how many people is it helping?
        - How easily can it be solved with existing tools like LLMs?
        - Are there any fairness and/or ethical issues of the tool?
        - How easy is it to maintain the tool and TCO?
        - Are there any liabilities or legal risk?
        - Worst case scenario harm if wrong (low, medium, high)?
        - Need to run local (y/n)?
- **How do we make decisions?**
    - The foundation makes decisions through a time boxed consensus mechanism. We are ruthlessly efficient for the greater good - Therefore, <u>decisions will be made by those who show up</u>.
        - With one exception: changing the mission, vision and principles require unanimous agreement among FAs.
    - The exceptions to the above rule are decisions that require Board approval. In those cases, a majority is required to approve decisions.
    - Foundation Advisor and Board of Directors decision boundaries:
        - Approval of annual budget and any unbudgeted expenditure exceeding $2500
        - Entering into contracts or legal agreements on behalf of the Foundation
        - Accepting grants or donations with conditions or restrictions
        - Appointing or removing Foundation Advisors
        - Appointing or removing Executive Committee members
        - Electing or removing directors and officers
        - Amending these Bylaws or the Articles of Organization
        - Authorizing legal action or settlements
        - Matters involving insurance, risk, and regulatory compliance
        - Dissolving the Foundation
    - Foundation Advisor decision boundaries (without Board involvement):
        - Idea pipeline
        - Technical Advisor and Contributor recruiting
        - Career coaching
    - Technical Advisor decision boundaries:
        - Technical direction
        - Roadmap and prioritization
        - Vision
        - Design, UI, UX
        - Products and tools
        - Idea pipeline
        - TA/Contributor recruiting
        - Definition of success and metrics
        - Evaluation methods/datasets,

        - TAs will facilitate these decisions through a consensus mechanism
        - Source code and commits
    - Collaborator decision boundaries:
        - PR merges
        - Releases
        - Refactors,
        - Technical implementation details
        - Contributor recruiting.
    - All Board of Director meeting minutes and decisions will be logged and made public via the foundation's GitHub page.
- **How do we connect projects with contributors and TAs?**
    - In phase-1, we have recruited and placed contributors to specific projects. Each TA owns one project - based on interest/skills.
    - In phase-2, we will look for contributors broadly, allow for self selection and manually intervene when needed. TAs will own 1-2 projects - based on interest/skills.
    - At scale we will use AI to match TAs and contributors to specific projects based on their background, skills, and interests. We will continue to allow for self selection at scale.
- **How do we find Foundation Advisors?**
    - Currently we are fully staffed in our advisory board. Find them here.
    - If we need more FAs will collectively recommend, reach out and onboard advisors or replace any advisor who vacates their position.
    - We will use the decision making framework above to determine who should be invited to serve as an advisor.
    - It is imperative to increase diversity of advisors based on gender, race, background, skills, geo location and other dimensions.
    - Over time we will grow Technical Advisors to become the next generation Foundation Advisors.
- **How do we find Technical Advisors?**
    - For phase-1 we are fully staffed from a TA perspective. Find them here.
    - For phase-2, we will use FAs networks, LinkedIn, and other mediums to identify potential TAs based on the profile described above.
    - Our approach for future phases is TBD.
- **How do we find contributors?**
    - In phase-1, we have used FA and TA networks and universities to identify people based on the profiles described above. Find them in project GitHub repos.
    - In phase-2 and beyond, we will expand our partnerships with universities and colleges, utilize their network and do broader outreach across different platforms, including social media.
- **How do we change b2's mission, vision, or principles?**
    - Per b2's Bylaws, the only way to change b2's mission, vision or principles is a unanimous vote among all FAs and Board of Directors of b2.
- **What tools do we use?**

    - GitHub as our code repo, issue tracking tool and backlog management mechanism. We will use GitHub to store and log decisions and other formal public documents. GitHub is our source of truth for everything that is public facing.
    - Google for working/internal docs, email system, decks, and spreadsheets.
    - Slack and gmail as our collaboration and communication channel.
- **How do we ensure openness and transparency?**
    - We open source all our code via GitHub and make it available to the public.
    - We publish decisions on our GitHub repo.
    - We identify FAs and TAs on our website (with opt-in). Their organizations will be used only as identification, not endorsement.
    - We publish all our code with the MIT or Apache-2.0 licenses.
    - Docs/content/decisions are published with CC BY or CC BY-SA licenses.
    - The foundation's name and logos are trademarks of The Benevolent Bandwidth Foundation. A detailed trademark policy will be created.
- **How does someone start contributing?**
    - Submit a collaboration form here - we guarantee a response!
    - Request and gain access to b2's GitHub, Google doc and slack environments
    - Find the project repo in GitHub
    - Fork the repo
    - Identify an issue to pick up, or create a new one
    - Execute and commit!
- **What about usage and adoption of AI we build?**
    - Tools without usage or adoption are not very useful. Plus, there is the cost of maintenance. Therefore, we will focus on promoting usage or adoption strategically if the tool isn't growing organically.
    - Sometime during Phase-1/Phase-2 we will be onboarding TAs and contributors directly focused on enabling, promoting and driving usage of the tools we deploy.
- **Do we maintain the tools we ship?**
    - We commit to approaching the tools we deploy as a product. We will continue to either invest in improving and maintaining them or we will deprecate them. We commit to not contributing to the sea of useless AI tools that rarely get used with stale code.